

# **Wildfire Detection through Image Analysis**

## **using CNN and MobileNetV2**




Artificial Intelligence & Machine Learning  
Alpen-Adria-University Klagenfurt

Evangelia Panteliadou

# Introduction

# Introduction

"Welcome to my presentation on wildfire detection using image analysis, where I'll be discussing how machine learning techniques can be leveraged to enhance early detection and response to wildfires."



# Motivation

- Personal Experience
- Environmental Impact
- Technological Potential




*Evia Island, Greece - Wildfire 2022 - Photo taken by Konstantinos Tsakalidis*

# Goals

- Primary Objective - an ML model that can accurately detect signs of wildfire
- Key Benefits
  - Timely Alerts - early warnings to local authorities
  - Damage Mitigation - improving preparedness and resource allocation
  - Technological Advancement - solving real-world environmental challenges with the use of AI

# Workflow





# Example of inputs - Fire and smoke class





# Example of inputs - No fire and nosmoke class





# Learning Task

# CNN Model

- Classification of images
  - **fire and/or smoke**
  - **no fire/no smoke**
- Train data
- Validate data
- Test data

# Tools and Libraries

- Data Preprocessing

- NumPy
- Pillow

- Development Environment

- Google Colab

- Model Design

- TensorFlow
- Keras
- MobileNetV2

- Data Analysis

- Pandas
- Matplotlib
- Seaborn

- Metrics

- scikit-learn

# Model Training



## # Model Training

```
checkpoint = ModelCheckpoint('models/best_model.keras', monitor='val_loss', save_best_only=True, verbose=1)
reduce_lr = ReduceLROnPlateau(monitor='val_loss', factor=0.2, patience=2, min_lr=0.00001)
early_stopping = EarlyStopping(monitor='val_loss', patience=3, restore_best_weights=True)

history = model.fit(
    train_generator,
    steps_per_epoch=steps_per_epoch,
    validation_data=validate_generator,
    validation_steps=validation_steps,
    epochs=5,
    callbacks=[checkpoint, reduce_lr, early_stopping],
    class_weight=class_weights
)
```

# Model Training - Results



```
Epoch 1/5
93/93 [=====] - ETA: 0s - loss: 0.1466 - accuracy: 0.9437
Epoch 1: val_loss improved from inf to 0.12397, saving model to models/best_model.keras
93/93 [=====] - 708s 8s/step - loss: 0.1466 - accuracy: 0.9437 - val_loss: 0.1240 - val_accuracy: 0.9643 - lr: 1.0000e-04
Epoch 2/5
93/93 [=====] - ETA: 0s - loss: 0.0881 - accuracy: 0.9700
Epoch 2: val_loss did not improve from 0.12397
93/93 [=====] - 249s 3s/step - loss: 0.0881 - accuracy: 0.9700 - val_loss: 0.1310 - val_accuracy: 0.9583 - lr: 1.0000e-04
Epoch 3/5
93/93 [=====] - ETA: 0s - loss: 0.0602 - accuracy: 0.9774
Epoch 3: val_loss improved from 0.12397 to 0.07729, saving model to models/best_model.keras
93/93 [=====] - 260s 3s/step - loss: 0.0602 - accuracy: 0.9774 - val_loss: 0.0773 - val_accuracy: 0.9821 - lr: 1.0000e-04
Epoch 4/5
93/93 [=====] - ETA: 0s - loss: 0.0521 - accuracy: 0.9838
Epoch 4: val_loss did not improve from 0.07729
93/93 [=====] - 259s 3s/step - loss: 0.0521 - accuracy: 0.9838 - val_loss: 0.0804 - val_accuracy: 0.9807 - lr: 1.0000e-04
Epoch 5/5
93/93 [=====] - ETA: 0s - loss: 0.0552 - accuracy: 0.9842
Epoch 5: val_loss improved from 0.07729 to 0.07381, saving model to models/best_model.keras
93/93 [=====] - 248s 3s/step - loss: 0.0552 - accuracy: 0.9842 - val_loss: 0.0738 - val_accuracy: 0.9792 - lr: 1.0000e-04
```

# Model Evaluation



```
# Model Evaluation
```

```
best_model = tf.keras.models.load_model('models/best_model.keras')
```

```
# Ensure that the test set is fully covered
```

```
test_steps = np.ceil(test_generator.samples / batch_size).astype(int)
```

```
train_generator.reset()
```

```
train_loss, train_accuracy = best_model.evaluate(train_generator, steps=train_generator.samples // batch_size)
```

```
print(f'Train Accuracy: {train_accuracy * 100:.2f}%')
```

```
validate_generator.reset()
```

```
val_loss, val_accuracy = best_model.evaluate(validate_generator, steps=validate_generator.samples // batch_size)
```

```
print(f'Validation Accuracy: {val_accuracy * 100:.2f}%')
```

```
test_generator.reset()
```

```
test_loss, test_accuracy = best_model.evaluate(test_generator, steps=test_generator.samples // batch_size)
```

```
print(f'Test Accuracy: {test_accuracy * 100:.2f}%')
```



# Model Evaluation - Output



```
93/93 [=====] - 199s 2s/step - loss: 0.0246 - accuracy: 0.9926  
Train Accuracy: 99.26%  
21/21 [=====] - 34s 2s/step - loss: 0.0732 - accuracy: 0.9792  
Validation Accuracy: 97.92%  
70/70 [=====] - 530s 8s/step - loss: 0.3665 - accuracy: 0.8987  
Test Accuracy: 89.87%
```

# **Preliminary Data Analysis**

# Data

- Data Sources
  - Combination of different Kaggle datasets
- Data Categories
  - images labeled
    - fire\_smoke
    - nofire\_nosmoke
- Data Characteristics
  - different times of day, weather conditions, and landscapes
- Visualization
  - initial examples of the data and distribution

# Data Augmentation and Image Preprocessing

- setting the image height and width for compatibility with MobileNetV2
- using a batch size of 32 for efficient processing
- data augmentation like rescale, rotation, etc. for diversity

```
# Image dimensions
img_height, img_width = 224, 224
batch_size = 32 # Increased batch size for faster processing

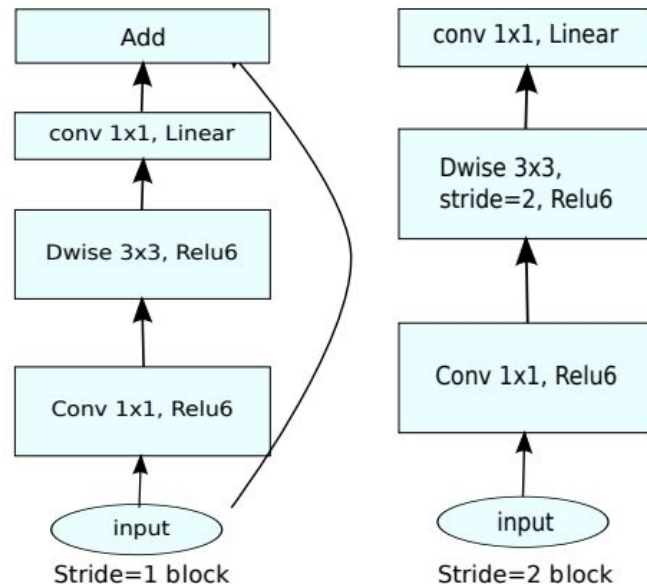
# Data Augmentation and Generators
train_datagen = ImageDataGenerator(
    rescale=1./255,
    rotation_range=20,
    width_shift_range=0.2,
    height_shift_range=0.2,
    shear_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True,
    fill_mode='nearest'
)

val_test_datagen = ImageDataGenerator(rescale=1./255)
```

# Implementation

# MobileNetV2 Classification - Architecture

MobileNetV2 is a classification model developed by Google. It provides real-time classification capabilities under computing constraints in devices like smartphones. This implementation leverages transfer learning from ImageNet to a dataset.



(d) Mobilenet V2



# Custom Layers

```
[16] # Load pre-trained MobileNetV2 model + higher level layers
base_model = tf.keras.applications.MobileNetV2(weights='imagenet', include_top=False, input_shape=(224, 224, 3))

# Add top layers
x = base_model.output
x = Flatten()(x)
x = Dense(64, activation='relu')(x)
x = Dropout(0.5)(x)
predictions = Dense(1, activation='sigmoid')(x)

# Create final model
model = Model(inputs=base_model.input, outputs=predictions)

# Freeze the base_model layers
for layer in base_model.layers:
    layer.trainable = False

# Compile the model
optimizer = Adam(learning_rate=0.0001) # Reduced learning rate
model.compile(optimizer=optimizer, loss='binary_crossentropy', metrics=['accuracy'])
model.summary()
```

# **Obtained Results and Analysis**

# Model Evaluation Metrics

71/71 [=====] - 118s 2s/step

Length of y\_true: 2250

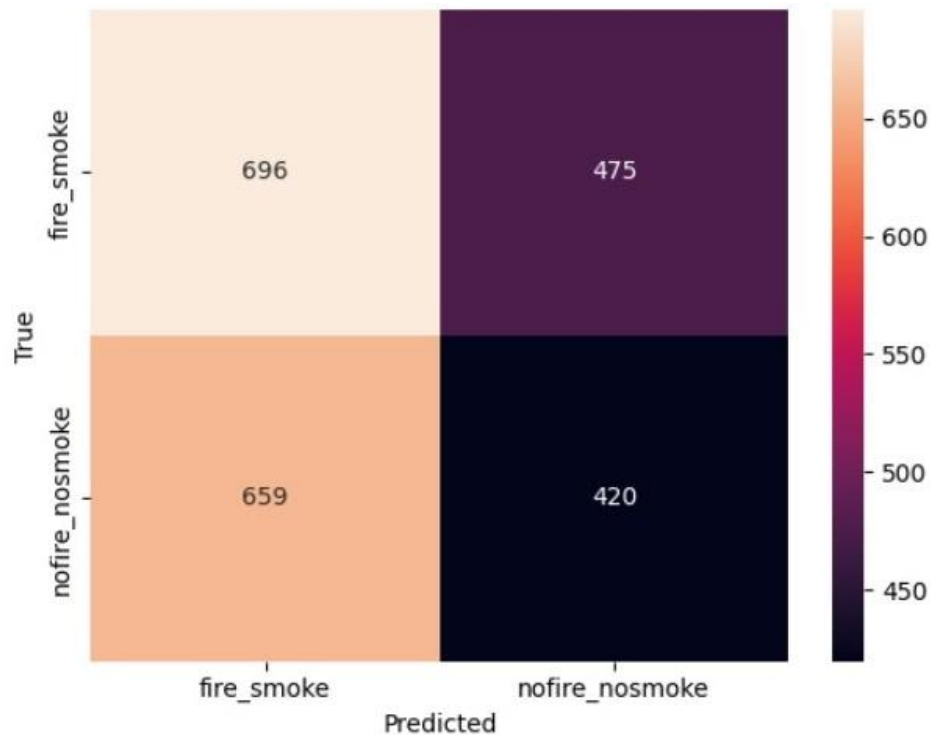
Length of y\_pred\_classes: 2250

Precision: 0.47



Recall: 0.39

F1 Score: 0.43

	precision	recall	f1-score	support
fire_smoke	0.51	0.59	0.55	1171
nofire_nosmoke	0.47	0.39	0.43	1079
accuracy			0.50	2250
macro avg	0.49	0.49	0.49	2250
weighted avg	0.49	0.50	0.49	2250



# Training and Validation Performance



**Thank you for your  
attention!**